

Introduction to Database Systems  
**Intro to Query Processing**

Literature: PDBM 13.1

Björn Þór Jónsson

# First Things First

- **All exercises now ONLINE on Discord!**
  - See post: <https://piazza.com/class/kdy92u8mvyh2s6?cid=29>
- Homework 3 due tomorrow
  - = DDL + Normalization + Indexing (+ SQL)
  - Question 2.1: See video: <https://use.vg/DtvIyrpjS7Zm>
- **Homework 2: Review next Monday at 18:00**
  - On the regular Zoom meeting!
- Now, let's finish last week's slides!

# Recap: Indexing

- Indexes are data structures that facilitate access to data from disk
  - ... if conditions are a prefix of indexed attributes
  - *Clustered* indexes store tuples that match a range condition together
  - Some queries can be answered looking **only** at the index (a *covering* index for query)
  - Indexes slow down updates and insertions
- The choice of whether to use an index is made by the DBMS *for every instance of a query*
  - May depend on query parameters
  - Don't have to name indexes when writing queries

# Today's Lecture

## Part A:

- Processing simple selections
- Processing complex selections
- Intro to join evaluation algorithms
- (Super brief) Intro to grouping / aggregations

## Part B:

- Short intro to database tuning
  - A very interesting and important topic!
  - Not necessary for the course, and will not get to it
  - I'll leave the slides in for those interested

# Part A: Query Processing

# Processing Simple Selections

- We discussed this earlier, to summarise:
- Point and range queries on the attribute(s) of the **clustered index** are almost always best performed using an index scan
- **Unclustered** indexes should only be used with **high selectivity queries**
- Exception: **Covering index** is good for any selectivity
- If no index exists, a full table scan is required!

# Processing Complex Selections

- We consider the conjunction ("and") of equality and range conditions.
- No relevant index: Full table scan
- One index relevant:
  - Highly selective: Use that index
  - If not: Full table scan
- Multiple relevant indexes:
  - One is highly selective: Use that index
  - No single condition matching an index is highly selective: Can "intersect" the returned sets

# Using a Highly Selective Index

- Basic idea:
  - Retrieve all matching tuples (few)
  - Filter according to remaining conditions
- If index is clustered or *covering*: Retrieving tuples is particularly efficient, and the index does not need to be highly selective.



# Using Several Less Selective Indexes

- For several conditions  $C_1, C_2, \dots$  matched by indexes:
  - Retrieve the addresses  $R_i$  of tuples matching  $C_i$ .
    - The addresses are in the index leaves!
  - Compute the intersection  $R = R_1 \cap R_2 \cap \dots$
  - Retrieve the tuples in  $R$  from disk (in sorted order)
- Remaining problem:
  - How can we estimate the selectivity of a condition?  
Of a combination of conditions?
  - Use some stats and probabilistic assumptions...

# Example

```
SELECT title
FROM Movie
WHERE year = 1990
      AND studioName = 'Disney' ;
```

## Examples of strategies:

1. Make a scan of the whole relation.
2. Find movies from 1990 using index, then filter.
3. Find Disney movies using index, then filter.
4. Combine *two* indexes to identify rows fulfilling both conditions.
5. Use *one* composite index to find Disney movies from 1990.
6. Find Disney movies from 1990 *and* their titles in a composite *covering* index.

# Example

```
SELECT title
FROM Movie
WHERE year = 1990
      AND studioName = 'Disney';
```

## Examples of strategies:

1. Make a scan of the whole relation.
2. Find movies from 1990 using index, then filter.
3. Find Disney movies using index, then filter.
4. Combine *two* indexes to identify rows fulfilling both conditions.
5. Use *one* composite index to find Disney movies from 1990.
6. Find Disney movies from 1990 *and* their titles in a composite *covering* index.

```
CREATE INDEX yearIdx
ON Movie(year)
```

```
CREATE INDEX studIdx
ON Movie(studioName)
```

```
CREATE INDEX yearStudIdx
ON Movie(year, studioName)
```

```
CREATE INDEX coveringIdx
ON Movie(year, studioName, title)
```

# Example – Variant 1

```
SELECT title
FROM Movie
WHERE year = 1990
      AND studioName = 'Disney';
```

## Examples of strategies:

1. Make a scan of the whole relation.
2. Find movies from 1990 using index, then filter.
3. Find Disney movies using index, then filter.
4. Combine *two* indexes to identify rows fulfilling both conditions.
5. Use *one* composite index to find Disney movies from 1990.
6. Find Disney movies from 1990 *and* their titles in a composite *covering* index.

## Available Indexes:

```
CREATE INDEX yearIdx
ON Movie(year)
```

```
CREATE INDEX studIdx
ON Movie(studioName)
```

Which strategies are possible and which index would be used?

# Example – Variant 2

```
SELECT title
FROM Movie
WHERE year = 1990
      AND studioName = 'Disney';
```

## Examples of strategies:

1. Make a scan of the whole relation.
2. Find movies from 1990 using index, then filter.
3. Find Disney movies using index, then filter.
4. Combine *two* indexes to identify rows fulfilling both conditions.
5. Use *one* composite index to find Disney movies from 1990.
6. Find Disney movies from 1990 *and* their titles in a composite *covering* index.

## Available Indexes:

```
CREATE INDEX yearIdx
ON Movie(year)
```

```
CREATE INDEX yearStudIdx
ON Movie(year, studioName)
```

Which strategies are possible and which index would be used?

# Example – Variant 3

```
SELECT title
FROM Movie
WHERE year = 1990
      AND studioName = 'Disney';
```

## Examples of strategies:

1. Make a scan of the whole relation.
2. Find movies from 1990 using index, then filter.
3. Find Disney movies using index, then filter.
4. Combine *two* indexes to identify rows fulfilling both conditions.
5. Use *one* composite index to find Disney movies from 1990.
6. Find Disney movies from 1990 *and* their titles in a composite *covering* index.

## Available Indexes:

```
CREATE INDEX yearIdx
ON Movie(year)
```

```
CREATE INDEX yearStudIdx
ON Movie(year, studioName)
```

```
CREATE INDEX coveringIdx
ON Movie(year, studioName, title)
```

Which strategies are possible and which index would be used?

# Example – Variant 4

```
SELECT title
FROM Movie
WHERE year = 1990
      AND studioName = 'Disney';
```

## Examples of strategies:

1. Make a scan of the whole relation.
2. Find movies from 1990 using index, then filter.
3. Find Disney movies using index, then filter.
4. Combine *two* indexes to identify rows fulfilling both conditions.
5. Use *one* composite index to find Disney movies from 1990.
6. Find Disney movies from 1990 *and* their titles in a composite *covering* index.

## Available Indexes:

```
CREATE INDEX idIdx
ON Movie(id)
```

```
CREATE INDEX titleIdx
ON Movie(title)
```

Which strategies are possible and which index would be used?

# Processing Complex Selections Revisited

- We have considered the conjunction (“and”) of a number of equality and range conditions.
- What about disjunctive (“or”) selections?
  - One full table scan  
OR
  - Multiple “and” queries



# Query Evaluation in a Nutshell

- SQL rewritten to (extended) relational algebra
- The building blocks in DBMS query evaluation are algorithms that implement relational algebra operations.
  - **Join** is the most important one!
- May be based on:
  - Reading everything / Sorting / Hashing
  - Using indexes can sometimes help!
- The DBMS optimizer knows the characteristics of each approach, and attempts to use the best one in a given setting

# Join Evaluation in a Nutshell

- Join is the most important operation!
- May be based on:
  - Reading everything / Sorting / Hashing
  - Using indexes can sometimes help!
- We consider a simple join:  
     $R \text{ JOIN } S \text{ ON } S.ID = R.ID$ 
  - Extends to more complex joins in a straightforward way

# Nested Loops Join

- The following basic algorithm can be used for **any** join:

```
for each tuple in R
  for each tuple in S
    if r.ID = s.ID
      then output (r, s)
```

- If the join condition is complex/broad, sometimes this is the only/best choice

R JOIN S ON S.ID <> R.ID

# Role of Index in Nested Loops Join

- If there is an index that matches the join condition, the following algorithm can be considered:
  - For each tuple in R
    - use the index to locate matching tuples in S
- Good if  $|R|$  is small compared to  $|S|$
- If many tuples match each tuple, a clustered or covering index is preferable.
- MySQL currently implements only this join algorithm and the previous naïve alternative.

# Example

```
SELECT *  
FROM Movie M, Producer P  
WHERE M.year=2015  
      AND P.birthdate<'1940-01-01'  
      AND M.producer = P.id;
```

## Some possible strategies:

1. Use index to find 2015 tuples, use index to find matching tuples in Producer.
2. Use index to find producers born before 1940, use index to find matching movies.
3. NL join Movie and Producer, then filter.

# Problem session

What would be good indexes for these queries?

1. 

```
SELECT firstNames
FROM person
WHERE gender='m'
      AND firstnames LIKE 'Maria%';
```
2. 

```
SELECT A.street, A.streetno
FROM person P
      JOIN address A ON A.person_id=P.id
WHERE P.lastname='Bohr'
      AND P.firstnames LIKE 'Niels%';
```


# Merge Join

Consider  $R \text{ JOIN } S \text{ on } R.ID = S.ID$

- Step 0: Sort  $R$  and  $S$  on  $ID$
- Step 1: Merge the sorted  $R$  and  $S$
- Cost:
  - If already sorted:  $O(|R| + |S|)$
  - Can we do better?
  - If not sorted:  
 $O(|R|\log|R| + |S|\log|S| + |R| + |S|)$

# Role of Indexes in Merge Joins

- Indexes can be used to read data in sorted order
- When is this a win?
  - Index is clustered
  - Index is covering
- When is this a loss?
  - Index is unclustered






# Hash Join

Consider  $R \text{ JOIN } S \text{ on } R.ID = S.ID$

- Best if  $S$  fits in RAM
- Step 0: Create a good hash function for ID
- Step 1: Create a hash table for  $S$  in memory
- Step 2: Scan  $R$  and look for matching tuples in the hash table
- Cost:  $O(|R| + |S|)$ 
  - Can we do better?
  - What if  $S$  does not fit in RAM?

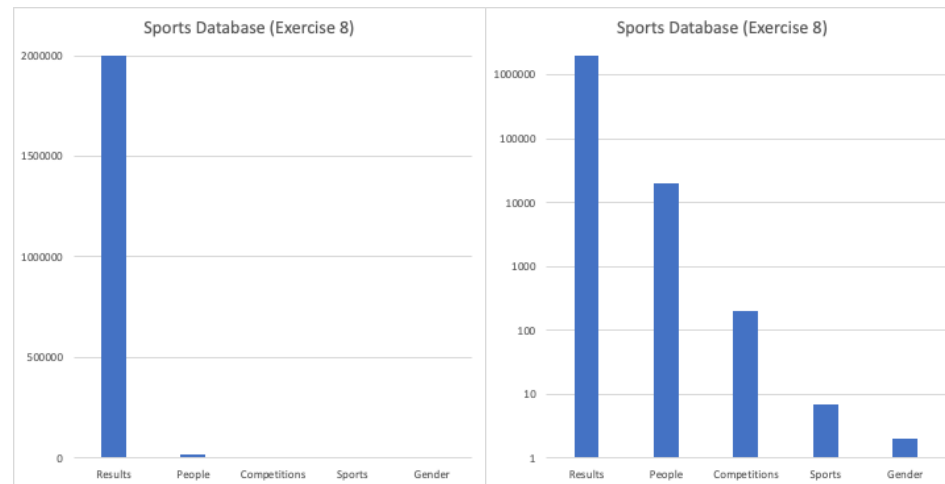
# Role of Indexes in Hash Joins

- Hash joins read all the relations
- How can indexes be useful?
  - Apply to non-join conditions
  - Index is covering



# Comparison of Join Algorithms

- Nested loops join:
  - Very costly  $O(|R| * |S|)$
  - Works for any condition → sometimes only option
- Merge join:
  - Works well if data is well clustered
  - Works well if relations are large and similar in size
- Hash join:
  - Works well if one relation is small
  - Is that often the case?



# Grouping Operations

- Many operations are based on grouping records (of one or more relations) according to the values of some attribute(s):
  - Join (group by join attributes, see above)
  - Group by and aggregation (obvious)
  - Set operations (group by all attributes)
  - Duplicate elimination (group by all attributes)
- Most database systems implement such grouping efficiently using sorting or hashing

# **Part B: Intro to Database Tuning**

# Database Tuning

- Done: ER design, database schema, views
  - Logical schema
  - External schema
- Next: Reach performance goals
  - Before throwing hardware at it...
  - ... consider the physical schema

## Two main tuning techniques:

- Adding indexes
- Changing the schema / physical storage

# Need to Understand the Workload!

- The major queries + frequency
  - Which tables and columns are read
  - Which columns appear in selections/joins?  
What is the likely reduction factor of the selections?
  - Why?  
Because indexes can speed up queries
    - Find records quickly !
- The major updates + frequency
  - As before +
  - Update type (UPDATE/INSERT/DELETE) and updated columns
  - Why updates?  
Because indexes can speed up OR slow down updates

# Query Optimization vs Query Tuning

- **Query optimization** is the process where the DBMS tries to find the "best possible" way of evaluating a given query.
  - Standard approach builds on finding a "good" relational algebra expression and then choosing how and in what order the operations are to be executed.
- **Query tuning** is a "manual" effort to make query execution faster.



# Query Opt vs Query Tuning

## Query Optimisation

- `SELECT a  
FROM R  
WHERE b IN (  
    SELECT b FROM S);`
- `SELECT DISTINCT a  
FROM R, S  
WHERE R.b = S.b;`

*From Lecture 3*

## Query Tuning

- `select R.peopleID, R.sportID, R.result  
from Results R  
where R.result = (  
    select max(R1.result)  
    from Results R1  
    where R1.sportID = R.sportID));`
- `select R.peopleID, R.sportID, R.result  
from Results R  
where (R.sportID, R.result) in (  
    select R1.sportID, max(R1.result)  
    from Results R1  
    group by R1.sportID));`

*From Homework 1*

# Indexing Decisions

- Starting point:
  - Keys (primary key / unique) are enforced using indexes
- Read / Write ratios
  - OLTP = 60 / 40  $\leftrightarrow$  OLAP = 98 / 2
- Which tables should be indexed?
- Should we create multiple indexes?
- Which columns should become search keys?
- Should the indexes be clustered or not?

# Choosing Columns

- Candidates for index search keys
  - Columns in WHERE clauses
  - Columns in GROUP BY clauses
  - Columns in ORDER BY clauses
- Columns that are rarely candidates
  - Large columns (too much space)
  - Frequently updated columns (too much maintenance)
  - Columns in SELECT clauses (not used to find tuples)
    - ... but see covering indices!

# Denormalization

- Normalization reduces redundancy and avoids anomalies
- Normalization can **improve** performance
  - Less redundancy => more rows/page => less I/O
  - Decomposition => more tables => more clustered indexes => smaller indexes
- The price of normalization:
  - Need to do more joins.

```
SELECT S.Name, T.Grade
FROM Student S, Transcript T
WHERE S.Id = T.StudId
      AND T.CrsCode = 'CS305'
      AND T.Semester = 'S2002'
```

# Denormalization

- **Tradeoff:** *Judiciously* introduce redundancy to improve performance of certain queries
- **Example:** Add attribute *Name* to Transcript

```
SELECT T.Name, T.Grade
FROM   TranscriptDenorm T
WHERE  T.CrsCode = 'CS305' AND T.Semester = 'S2002'
```

- Join is avoided
- If queries are asked more frequently than Transcript is modified, added redundancy might improve average performance
- But, Transcript' is no longer in BCNF since key is (*StudId*, *CrsCode*, *Semester*) and *StudId* → *Name*

# Problem Session

- Schema:
  - Customer(cno, name, country, type)
  - Invoice(ino, cno, amount)
  - Additional indexes on country and amount.
- What are possible query plans for this?
- Can denormalization help?

```
SELECT ino, amount
FROM customer, invoice
WHERE country="Sweden"
      AND amount > 10000
      AND customer.cno = invoice.cno
```

# Denormalization

- Denormalized schema:
  - Customer(cno, name, country, type)
  - Invoice(ino, cno, amount, **country**)  
redundant attribute

- Query can run on the new Invoice:

```
SELECT ino, amount
FROM invoice
WHERE country="Sweden"
      AND amount > 10000
```

- What speaks against such denormalization?

# Materialized/Indexed Views

- Some systems (including PostgreSQL) allow storing (materializing) views as tables
  - System keeps the view up to date
  - System automatically chooses to use the view

```
CREATE MATERIALIZED VIEW StudTransJoin  
AS
```

```
SELECT S.Name, T.Grade  
FROM   Student S, Transcript T  
WHERE  S.Id = T.StudId
```

```
SELECT S.Name, T.Grade  
FROM   Student S, Transcript T  
WHERE  S.Id = T.StudId  
       AND T.CrsCode = 'CS305'  
       AND T.Semester = 'S2002'
```



# Partitioning of Tables

- A table might be a performance bottleneck
  - If it is heavily used, causing locking contention (more on this later in course)
  - If its index is deep (table has many rows or search key is wide), increasing I/O
  - If rows are wide, increasing I/O
- Table partitioning might be a solution to this problem.

# Horizontal Partitioning

- If accesses are confined to disjoint subsets of rows, partition table into smaller tables containing the subsets
  - Geographically, organizationally, active/inactive
- Advantages:
  - Spreads users out and reduces contention
  - Rows in a typical result set are concentrated in fewer pages
- Disadvantages:
  - Added complexity
  - Difficult to handle queries over all tables

# Vertical Partitioning

- Split columns into two subsets, replicate key
- Useful when table has many columns and
  - it is possible to distinguish between frequently and infrequently accessed columns
  - different queries use different subsets of columns
- **Example:** Employee table
  - Columns related to compensation (tax, benefits, salary) split from columns related to job (department, projects, skills).
- DBMS trend (for analytics):
  - **Column stores**, with *full* vertical partitioning.
  - More on this next week.

# Take Aways

- Performance difference between well-tuned and poorly-tuned applications can be massive!
- The DBMS does its best to optimize queries, but sometimes it needs help!
  - Query tuning – rewrite as joins or non-correlated subqueries
  - Indexes – solve 90+% of all other performance problems
- If that is not sufficient:
  - Materialized views / Partitioning / Denormalization
  - Beyond the scope of this course!